

# Revamping JavaScript Static Analysis via Localization and Remediation of Root Causes of Imprecision

## ABSTRACT

The effectiveness of static analysis is challenged by the dynamic features of programming languages such as JavaScript, often resulting in impractical results in terms of performance and precision for specific programs. It is a time-consuming but crucial process during static analysis design to discover the sources of analysis impracticality. We present a systematic support for root-cause localization that focuses on JavaScript points-to analysis. When inspecting the intermediate results, the history information of points-to propagation, obtained through a labeled propagation system, is queried to identify the program pointers that have a big impact on overall analysis performance and/or precision. We also have designed an improvement suggestion algorithm that takes dynamic information as inputs to simulate the benefits of various context-sensitive analyses and recommends the appropriate context sensitivity for the root cause functions. The experimental results on JavaScript library applications show that our root-cause localization algorithm can accurately identify the specific program constructs that render existing static analysis impractical. We also have shown that static analysis applying the context sensitivity based on our improvement suggestion algorithm achieves a good balance between precision and performance for the benchmark programs.

## 1. INTRODUCTION

Dynamic programming languages such as JavaScript are now in widespread use for both client-side and server-side applications, often together with cloud services and/or mobile devices. The flexibility of these languages, for example for building prototypes, is key to their popularity, but their dynamic nature presents real challenges to static program analyses. These challenges affect analyses used to ensure the security of applications (e.g., integrity and confidentiality of data), to optimize code for good performance and to aid maintenance programmer understanding of code in need of updating. For example, JavaScript allows object prop-

erty accesses as associative arrays where the property name may be set as a result of execution. Making assumptions more accurate than worst-case in this case may be difficult. Moreover, JavaScript supports multiple programming paradigms including object-oriented, functional and procedural programming. Each programming paradigm requires specialized techniques for accurate analysis. These are examples of how JavaScript properties render accurate static analysis very difficult for many real-world programs.

The effectiveness of a static analysis often is evaluated as a combination of its precision (i.e., low false positive rate) and performance (i.e., efficiency in a limited time/space budget). Often with complex, medium-sized JavaScript programs, especially those using real-world libraries such as *jQuery*, static analysis cannot reach a useful solution within a reasonable time budget despite recent progress on improving the state-of-the-art (e.g., [13, 2, 15, 7]). Often it is even difficult to obtain a good approximation of the program call graph.

We present a new approach to tackle this problem in JavaScript programs. When applying an impractical static call graph construction and points-to analysis algorithm to the program, extra information about points-to propagation is gathered in the points-to graph under construction. Note that the accuracy of the points-to graph of the objects at call sites directly determines the accuracy of the call graph. A heuristic process observes the analysis propagation phase in order to catch anomalous behavior (i.e., when the analysis is becoming too approximate through propagation of inaccuracy). A diagnosis algorithm is applied to trace this “bad” behavior back to its root causes, which when found, present function calls for which judicious application of specific context sensitivity policies will circumvent the anomalous behavior after analysis restart. This process utilizes dynamic analysis results in addition to the static analysis self-inspection to help choose the kinds of context sensitivity to propose. We call this entire process *root-cause localization* and it is crucial for designing effective, new static analysis algorithms for JavaScript.

More specifically, the systematic support for root-cause localization focuses on points-to analysis, an enabling static analysis for various automated software tools. For an unscalable and/or too imprecise static points-to analysis on a target program, we keep track of the history information in the propagation system, labeling the origins of the points-to relations. By examining the analysis propagation as it is performed, heuristics can be applied to decide when to perform root-cause localization (i.e., when the analysis result is be-

coming anomalous). Our automatic root-cause localization uses the intermediate points-to results and the additional labels to identify the variables and/or properties that have a big impact on overall analysis precision and performance as the *root causes*. In addition, we have designed an automatic approach that suggests specialized techniques for improving analysis precision at these critical program constructs. Using a dynamic trace of program execution, we build a set of dynamic points-to graphs using various kinds of context sensitivity; the idea is that the dynamic points-to graphs can simulate the possible effects of a particular kind of context sensitivity to be used at the root cause function in the re-started analysis.

This automatic root-cause localization relieves a static analysis designer from the chores of manually inspecting the program and the analysis implementation to understand the sources of imprecision. Moreover, the automatic improvement suggestion provides possible context-sensitive analysis choices that may significantly improve the overall analysis performance and precision. The specialized analysis configurations derived from the results of our approach, with possible adjustment from the static analysis designer, can be executed on the same program to observe if the performance and/or precision issues have been resolved. If necessary, the same process can be iteratively performed to locate sufficiently many of the sources of imprecision in the analysis for the target program to achieve the desired accuracy and performance. In support of this approach, we have conducted experiments to evaluate the accuracy of the automatic root-cause location and the subsequent improvement suggestions on real-world JavaScript libraries and applications.

The major contributions of this work are:

- We present the *first research* that focuses on supporting static analysis design with automatic root-cause localization, identifying the sources of imprecision that have a big impact on analysis precision and performance.
- Our approach is the *first* to use dynamic information to automatically suggest the kind of context sensitivity needed for significant precision improvement on identified root causes of analysis inaccuracy.
- We present an evaluation of the proposed approaches on several benchmarks. The experimental results on JavaScript library applications demonstrate that our root-cause localization algorithm accurately identifies the program constructs that cause an initial static analysis not finishing in a 10 minute time allotment, applying specialized context sensitivity on which significantly improves the analysis performance. The results on JavaScript benchmarks also show a context-sensitive analysis that selectively applies the recommended context sensitivity from our automatic improvement suggestion achieves a much better balance between precision and performance compared to both an imprecise and an expensive analysis.

In Section 2, we motivate this project using empirical results and an example. Section 3 presents an overview of the process. Sections 4 and 5 describe the details of root-cause localization and improvement suggestion algorithms, respectively. Section 6 presents experimental results. Section 7 discusses related work and Section 8 offers our conclusions.

## 2. BACKGROUND & MOTIVATION

In this section, we first introduce the background on JavaScript points-to analysis with an example. We then discuss an empirical study of interesting static analysis behavior that guided our design. Finally, we intuitively discuss our approach using the code example.

### 2.1 Background

Points-to analysis approximates the program’s heap by calculating the set of abstract values a variable or reference property may have during execution. Context sensitivity is a general technique to achieve more precise program analysis by distinguishing between calls to a function [9]. It has been demonstrated that applying specialized context sensitivity in JavaScript points-to analysis is an effective approach for improving its precision and performance. For example, Sridharan et al. [13] apply the values of a parameter *p* of a function as the calling context, if *p* is used as the property name in a property access (e.g., *v[p]*), as a special treatment for dynamic property accesses in JavaScript. Because this algorithm is designed to address the challenges caused by a specific language feature, despite the fact that it performs well for the programs where this feature is present, other program constructs found in real-world JavaScript applications may render the analysis ineffective, which require more accurate handling from the static analysis.

```

1 function extendBasic(target, source) {
2     var name;
3     target = target || {};
4     for (name in source) {
5         target[name] = source[name];
6     }
7     return target;
8 }
```

Figure 1: Modified `extend` function of jQuery 1.6.1.


Therefore, an important stage of static analysis design is to identify the causes of a static analysis producing unexpected results. Intuitively, we define a *root cause* as a program construct that is a source of the precision and/or performance loss for a static analysis. Specifically, if the overall precision and performance of an analysis *A* improves significantly via a specialized handling of the program construct in a specific program location, this construct is the root cause of impracticality of the analysis *A* for the program. For example, Figure 1 shows a small piece of code from *jQuery*, the most widely used JavaScript library [14]. A whole-program 1-CFA analysis [10] that separately analyzes each different call site of a function has scalability problems for any simple application that uses *jQuery* [13]. Applying the technique proposed by Sridharan et al. [13] only to the property accesses at line 5 resolves its performance issues. Therefore, this program construct is a root cause of the 1-CFA analysis’ impracticality for *jQuery* applications.

Unfortunately, identifying the root causes is a high cost process, requiring extensive experience in designing static analyses as well as deep understanding of the target programs. To the best of our knowledge, there is little tool support for this process, making it time-consuming and hard to be comprehensive. For example, identifying the property accesses at line 5 in Figure 1 as the root cause for 1-CFA

analysis is difficult because (i) the *jQuery* library consists of about 9,000 lines of code, and (ii) similar program constructs are used throughout the *jQuery* library, whereas this particular program location is critical to the overall analysis results. Therefore, we are motivated to develop new techniques to assist in *root-cause localization* by automating part of the process of identifying the sources of impracticality.

## 2.2 Static Analysis Behavior

We have performed a brief empirical study to understand the behavior of JavaScript static analysis. Figures 2 and 3 show the different behaviors of two points-to analyses of a simple application that uses *jQuery*. The two points-to analyses in comparison are the 0-1-CFA analysis (i.e., only uses 1-CFA analysis for the constructors to name objects by their allocation sites) and the combined 1-CFA and specialized argument-sensitive analysis [13]. Overall, the 0-1-CFA analysis experiences performance and precision issues (e.g., the analysis cannot finish analyzing the program within a time budget of 10 minutes), while the combined context-sensitive analysis performs significantly better.

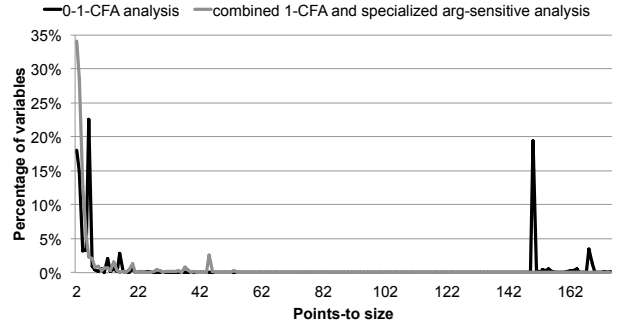


**Figure 2:** Points-to size growth during the analysis lifetime.

Figure 2 shows the trend of overall points-to size (i.e., total number of points-to edges) growth for these two analyses during their lifetimes. The x axis presents the number of evaluations<sup>1</sup> and y axis presents the total points-to size of all variables in the program. The good combined context-sensitive analysis’ points-to size grows steadily “linear” throughout its lifetime. On the other hand, the overall points-to size growth of 0-1-CFA analysis exhibits “jumps”, the periods during which its overall points-to size dramatically increases. For example, between evaluations of 23,000 and 25,000, the overall points-to size of the 0-1-CFA analysis grows about ten times. The existence of such “jumps” indicates that the overly-approximated results are frequently propagated, resulting in significantly overall precision loss. In addition, since the 0-1-CFA analysis experiences scalability issues, it remains incomplete at the end of the allocated analysis time and its overall points-to size keeps growing after the 120,000 evaluations in Figure 2.

Figure 3 shows the distributions of the points-to sizes for each variable in the program. The x axis presents the points-to size of a variable and the y axis presents the percentage of the variables in the program with the corresponding points-to size. For the combined context-sensitive analysis, the points-to size of the majority of the variables (i.e., 81%)

<sup>1</sup>An evaluation in the points-to analysis solves a constraint that may result in changes of the points-to results.



**Figure 3:** Points-to size distribution.

is less than 6 and few variables are associated with large points-to sets. For the 0-1-CFA analysis, the points-to sizes of only 40% variables are less than 6 and there are condensed occurrences of variables with extremely large points-to sets (e.g., about 20% of the variables’ points-to sizes are 150). This result also indicates that overly-approximated results of specific program constructs may pollute many places in the program due to copying during the points-to propagations.

The above results demonstrate the significantly different behaviors between points-to analyses when their performance and precision vary. These results motivated us to design an automated approach to identify root causes of imprecision via the differences of static analysis behavior.


## 2.3 Root-cause Localization & Improvement Suggestion

We now use the example in Figure 1 to illustrate the ideas on localizing root causes. First, the root-cause localization should be performed during the period in which overall precision of the points-to analysis starts to decrease, reflected as the “jumps” in terms of the overall points-to size in Figure 2. Second, we use the history information of points-to propagations and the incomplete points-to results to locate the program constructs that are root causes of imprecision. Intuitively, two conditions should be met: (i) the program construct has a wide reach within the propagation system (e.g., the values of the property access `source[name]` are assigned to property `name` of `target` at line 5 and are transitively propagated to about 500 other program variables or reference properties), and (ii) the impact of its wide reach is significant (e.g., looking up the property `name` of `source` at line 5 produces the points-to size of 150). Therefore, the imprecision of this property access results in the overall imprecision of the 0-1-CFA analysis when analyzing *jQuery* applications, becoming a root cause of imprecision.

In addition, to further assist in the process of improving the results of static analysis, we design an improvement suggestion algorithm that uses dynamic information to suggest appropriate context sensitivity to improve the analysis precision on the identified root causes. The dynamic information is used to simulate the benefits of different context-sensitive analyses, a generally applicable idea to quantify the potential precision of a specific context sensitivity.

## 3. TECHNICAL OVERVIEW

Figure 4 summarizes the root-cause localization process with the support of our automated techniques. We discuss



**Figure 4:** Root-cause localization process overview.

the steps comprising the process in the following.

We first run the static analysis in its initial context sensitivity configuration, which may lead to performance and/or precision issues. To monitor the behavior of the analysis, we run it in *diagnostic* mode by instrumenting the propagation system with labels, the result being a *labeled propagation system* that keeps track of the history of point-to propagations by labeling the origins of points-to relations. We avoid from the common situation that the analysis crashes (i.e., running out of time/space budgets), and so the metadata for root-cause localization is lost, by obtaining snapshots of the analysis’ state periodically.

Later, in Section 4, we describe and motivate our heuristics to determine whether the observed propagation behavior is anomalous. Intuitively, an anomaly is flagged if a given variable or reference property, collectively referred to as *program pointers*, is associated with a large points-to set, and the label corresponding to the pointer propagates to many other pointers within a small number of evaluations of the propagation system. Here is where the labels come into play, in recording the transitive propagation of abstract objects across pointers. Our localization algorithm ranks the results by order of their impact on precision, thereby reflecting how likely they are to be among the root causes for imprecision.

The optional remediation step, where suggestions for improvement are provided, accepts as input (i) the candidate root causes, (ii) a set of context sensitivity policies, as well as (iii) one or more dynamic execution traces of the program. The significance of the dynamic traces is in providing fully precise points-to information, such that the effect of different sensitivity policies can be evaluated.

Intuitively, the remediation process maps between the root causes identified statically and the concrete points in the trace  $t$ , and simulates the effect of different sensitivity policies atop  $t$ . Given for example statement  $\ell: y = x[p]$  in method  $m$  as the root cause, the first step is to simulate the least precise treatment of  $\ell$  by merging together concrete points-to information from all its occurrences in  $t$ . What follows is an iterative process, wherein different context sensitivity, and combinations thereof, are applied, and their effects are simulated. A suggestion is output for  $\ell$  to utilize a certain combination  $C$  of context sensitivity if (i)  $C$  partitions the points-to set of  $x[p]$  effectively, and (ii) more involved combinations are only negligibly better.

The final step, following the automated techniques, is for the user (i.e., either the analysis designer or an end user capable of configuring the analysis) to manually adjust the analysis configuration based on the localization results and/or the improvement suggestions. Upon doing so, the analysis can be rerun under the adjusted configuration to observe if the performance and/or precision issues have been resolved. The same process can be performed iteratively to

locate all the root causes, thereby leading to a specialized analysis configuration that meets the performance/precision requirements of the user at hand.

In the next two sections, we dive into the details of the two main algorithms: root-cause localization and improvement suggestion. We discuss each in turn.

---

#### Proc 1 Root-cause localization workflow.

---

**Input:** `config`: analysis configuration

**Input:** `i`: evaluation interval

**Output:** `R`: set of root causes

```

1: sys  $\leftarrow$  initialize propagation system with config
2: while (c  $\leftarrow$  sys.<next constraint>)  $\neq$  NULL do
3:   for each assignment v1 = v2 from c do
4:     ptsv1 = ptsv1  $\cup$  ptsv2 //points-to propagation
5:     lv1 = lv1  $\cup$  lv2  $\cup$  {v2} //label propagation
6:   end for
7:   if (k  $\leftarrow$  # of evaluations) mod i = 0 then
8:     grow  $\leftarrow$  points-to size growth in past i evaluations
9:     if grow > threshold then
10:      g  $\leftarrow$  intermediate static labelled points-to graph
11:      for each pointer node n in g do
12:        impactn  $\leftarrow$  compute the impact of n on g
13:      end for
14:      R  $\leftarrow$  high impact nodes
15:      return
16:    end if
17:  end if
18: end while

```

---

## 4. ROOT-CAUSE LOCALIZATION

We first explain, more technically, how the labeled propagation system is implemented. We then move to a description of our indicators, based on the labels, whether an anomalous propagation behavior is taking place. The overall workflow of root-cause localization is shown in Procedure 1.

### 4.1 Labeled Propagation System

A propagation system for the points-to analysis solves the constraints to reach a fixpoint, propagating the points-to relations of the variables and reference properties in the program. A majority of the constraints that exist in the propagation system are assignments. For example, to process an invoke instruction, the generated constraints include assignments from the actual arguments to the formal parameters, from the callee’s return values to the left-hand side variable of the invoke instruction, etc.

We assume a subset-based (aka inclusion-based) propagation system [1], which means that the system solves a constraint that assigns a program pointer `v1` to another pointer `v2` by adding the points-to set of `v1` to that of `v2`. In a standard propagation system, there is no provenance. The points-to set associated with a pointer may be the result of direct or transitive assignments, and there is no telling in general how it evolved to its current state. Such information is critical, however, to identify root causes, since frequent assignments involving an inaccurate points-to set may pollute the overall precision of the points-to solution as depicted in Figures 2 and 3.

To address this loss of information, in our labeled propagation system, each constraint that assigns the values of

v2 to v1 results in not only the changes to v1's points-to relations but also a label v2 associated with the points-to set of v1, indicating that (some of) the points-to relations of v1 were propagated through v2 (lines 3-6 in Procedure 1). As an example, in WALA<sup>2</sup> intermediate representation (IR), which is in SSA form [3], the statement at line 5 in Figure 1 reduces to (a)  $v_{tmp} = source[name]$  and (b)  $target[name] = v_{tmp}$ , where  $v_{tmp}$ ,  $source$ ,  $target$  and  $name$  are all local variables of the function.

For the property read instruction (a), the analysis would first query the points-to set of  $source$ ,  $P_{source}$ , and the points-to set of  $name$ ,  $P_{name}$ . The pairs of each element in  $P_{source}$  and  $P_{name}$  (e.g.,  $P_{source_i} \cdot P_{name_j}$ ) are returned as the results of looking up the reference properties of  $source[name]$ . Note that the values of  $name$  iterate over all the property names of  $source$ , which in practice are the large set of function names loaded in *jQuery*. This ultimately results in a large points-to set for  $v_{tmp}$  due to multiple assignments from various reference properties. We retain all these reference properties (e.g.,  $P_{source_i} \cdot P_{name_j}$ ) as labels attached to the points-to set of  $v_{tmp}$ .

For the property write instruction (b), the analysis similarly looks up the reference properties of  $target[name]$  (e.g.,  $P_{target_i} \cdot P_{name_j}$ ) and adds the points-to relations of  $v_{tmp}$  to each of these reference properties, resulting in an overly approximated points-to set for each function name in *jQuery*. We retain both  $v_{tmp}$  and the existing transitive labels from  $v_{tmp}$  (e.g.,  $P_{source_i} \cdot P_{name_j}$ ) in the points-to set of each reference property of  $target[name]$ .

In addition to tracking the set of labels that are immediately or transitively propagated to the points-to set of a given pointer, we generate a *propagation-history graph* for the points-to set of that pointer, which organizes the points-to propagation history into a hierarchical structure. Intuitively, a node in the propagation-history graph is a program pointer that has bearing on the specific points-to set. An edge from node  $n_1$  to another node  $n_2$  represents that the points-to set of  $n_2$  was explicitly added to that of  $n_1$ . The entry points of the graph are the program pointers whose points-to sets are directly propagated into the corresponding points-to set.

For the same example in Figure 1, the propagation-history graph for the points-to set of  $v_{tmp}$  consists of all the reference properties of  $source[name]$  as entry points.  $v_{tmp}$  is the entry point of the propagation-history graph for the points-to set of each reference property of  $target[name]$  (e.g.,  $P_{target_i} \cdot P_{name_j}$ ), while there also are edges from  $v_{tmp}$  to the properties of  $source[name]$ .

As motivated above, during the propagation process, the labeled system is paused regularly upon completing cycles of  $i$  evaluations, for a fixed value of  $i$  (line 7 in Procedure 1). We base our analysis of root causes of impracticality on the resulting intermediate states.

Specifically, we count the total number of points-to relations (i.e., edges) in the intermediate points-to graph for the  $k_{th}$  pause (i.e., after  $n \times k$  evaluations),  $actual(k)$ . We then compare the value of  $actual(k)$  with the value of  $predict(k)$  to decide whether to perform root-cause localization.

For meaningful comparison, we utilize the results from the previous  $k-1$  pauses in performing linear regression analysis




Figure 5: Prediction via linear regression.

to find the fitted line  $y = \text{intercept} + \text{slope} \times x$ , where  $x$  is the number of evaluations and  $y$  is the total number of points-to relations. The number of points-to relations of the  $k_{th}$  pause can then be predicted:  $predict(k) = \text{intercept} + \text{slope} \times kn$ . If  $actual(k) > 110\% \times predict(k)$ , then we decide to perform the root-cause localization at the  $k_{th}$  pause; otherwise, we continue the points-to analysis under the labeled propagation system.

Figure 5 shows this prediction model when running the 0-1-CFA analysis on a *jQuery* application. The analysis is paused every 1000 evaluations and the fitted line in Figure 5 is calculated from the results of the first 23 pauses (i.e., 1000 to 23,000 evaluations). The predicted total points-to edges at the 24,000 evaluations is 16,234, while  $actual(24)$  is 88,857, significantly passing the threshold to perform root-cause localization. This prediction model can capture the “jump” phase of the points-to analysis shown in Figure 2, which is important to localize the root causes accurately.

If localization is performed too early, the over-approximate results may not have surfaced yet. If localization is performed late during the analysis, then the over-approximate results due to the root causes may have polluted a large portion of the overall points-to results, making it difficult to identify the root causes of imprecision. Moreover, the evaluation interval to pause the analysis,  $i$ , as well as the decision threshold, may be tuned based on the budget and goals of root-cause localization.

## 4.2 Identifying Root Causes for Divergence

When performing root-cause localization, we use the intermediate points-to graph and the associated labels to identify possible root causes (line 10 in Procedure 1). For each variable or reference property  $v$ , we count (i) its points-to size,  $|P_v|$ , as well as (ii) its number of occurrences as labels in the points-to sets of other variables and/or reference properties,  $|L_v|$ .  $|L_v|$  measures how widely  $v$  reaches within the propagation system, and  $|P_v|$  measures if the impact of its wide reach is significant. We therefore use the scoring heuristic  $S_v = |P_v| \times |L_v|$  as a measurement of the possibility of  $v$  being the root cause (lines 11-13 in Procedure 1).

The root-cause localization stage reports a set of variables as root causes in descending order of their scores. For example,  $v_{tmp}$ , the result of the property read instruction at line 5 in Figure 1, achieves an extremely high score (73,950) with  $|P_{v_{tmp}}| = 150$  and  $|L_{v_{tmp}}| = 493$  when localization is performed after 24,000 evaluations. Its score is 75-times higher

<sup>2</sup>[http://wala.sourceforge.net/wiki/index.php/Main\\_Page](http://wala.sourceforge.net/wiki/index.php/Main_Page)

than the variable with the second highest score (990), rendering it the only root-cause candidate for imprecision of the 0-1-CFA analysis for this *jQuery* application. We report the variables and/or reference properties whose scores are at least half the highest score as *candidate (or suspicious) root causes*.

---

**Proc 2** Improvement suggestion workflow.

---

**Input:**  $R = \{r_1, \dots, r_n\}$ : root causes

**Input:**  $t$ : dynamic trace

**Input:**  $CS$ : context sensitivity policies

**Output:**  $\langle V, S \rangle = \{(r_1, s_1), \dots, (r_n, s_n)\}$ : suggestions

```

1: for each  $cs \in CS$  do
2:    $G = G \cup \{gen(t, cs)\}$ 
3: end for
4: for each  $r \in R$  do
5:    $A_{\langle r, G \rangle} \leftarrow \emptyset$ 
6:   for each  $g \in G$  do
7:      $a_{\langle r, g \rangle} \leftarrow$  query  $r$ 's points-to size in  $g$  and measure
       its accuracy corresponding to  $g$ 's context sensitivity
8:      $A_{\langle r, G \rangle} = A_{\langle r, G \rangle} \cup \{a_{\langle r, g \rangle}\}$ 
9:   end for
10:   $a_{\langle r, g_{min} \rangle} \leftarrow \min(A_{\langle r, G \rangle})$ 
11:   $\langle V, S \rangle = \langle V, S \rangle \cup \{ \langle r, g_{min}.cs \rangle \}$ 
12: end for
```

---

## 5. IMPROVEMENT SUGGESTION

Building on the diagnostic algorithm that detects root causes, the next step is to automatically compute suggestions how to improve the analysis per the program, configuration and context sensitivity at hand. This is the focus of this section, whose workflow is shown in Procedure 2.

First, the target program is executed to collect a dynamic trace, recording the following run-time artifacts in order of occurrence: (i) function entries and exits; (ii) invocations; and (iii) property reads and writes. At a property read/write instruction, we record (i) the instruction location in the program; (ii) the allocation site of the base object (as its program location); (iii) the property name, and (iv) the allocation site of the value, if it is a reference object, or the type of the value, if it is a primitive value. At an invoke instruction, we record (i) the location of the call site; (ii) the location of the target function; (iii) the allocation site of the receiver object; and (iv) the allocation sites and/or the types of the actual arguments.

Second, dynamic points-to graphs based on the dynamic trace are generated per the available types of context sensitivity. Procedure 3 captures the algorithm that produces a dynamic points-to graph with respect to a specific context-sensitive analysis (e.g., 1-CFA, 1st-argument-sensitive [15], or context-insensitive analysis). The inputs are the dynamic trace, **trace**, and the type of context sensitivity, **cs**. The algorithm iterates through all the instructions recorded in the dynamic trace. For each instruction  $i$ , the algorithm examines its kind. If it is an invoke instruction, at line 5, the call site and the argument are pushed into the call stack, **stack**. If it exits a function, then top element of **stack** is removed at line 7. If it is a property read or write instruction, at lines 9 to 15, depending on the input context sensitivity, the calling context, **context**, is determined by the call site and the argument from the top element of **stack** for 1-CFA and

argument-sensitive analysis, respectively; the calling context is *everywhere* for context-insensitive analysis. In the dynamic points-to graph, a variable is represented by (i) the location of the instruction, (ii) the part of the instruction (i.e., base, property or value), and (iii) the calling context. At lines 16 to 18, the object allocation site, represented by program location, of each part of the instruction collected at runtime is assigned to the corresponding variable node in the dynamic points-to graph. Note that this algorithm is general in that various dynamic points-to graphs that can be generated under different context sensitivity.

---

**Proc 3** Dynamic points-to graph generation,  $gen(t, cs)$ .

---

**Input:**  $t$ : dynamic trace

**Input:**  $cs$ : context sensitivity

**Output:**  $g$ : dynamic points-to graph

```

1:  $stack \leftarrow \emptyset$ 
2: while ( $i \leftarrow next(t)$ )  $\neq$  NULL do
3:   switch (kindOf  $i$ )
4:     case INVOKE:
5:        $stack.push$ (call site and 1st arg of  $i$ )
6:     case FEXIT:
7:        $stack.pop$ 
8:     case PREAD || PWRITE:
9:       if  $cs = 1\text{-CFA}$  then
10:         $context \leftarrow$  immediate call site on  $stack$ 
11:       else if  $cs = 1st\text{-argument-sens}$  then
12:         $context \leftarrow$  immediate argument on  $stack$ 
13:       else if  $cs = context\text{-insens}$  then
14:         $context \leftarrow everywhere$ 
15:       end if
16:        $g(i_{loc}, base, context) \rightarrow i_{base}$ 
17:        $g(i_{loc}, property, context) \rightarrow i_{property}$ 
18:        $g(i_{loc}, value, context) \rightarrow i_{value}$ 
19:     end switch
20: end while
```

---

Now that we have obtained the dynamic points-to graphs under different context sensitivity policies (i.e., 1-CFA, argument sensitivity or context insensitivity), we can determine which of the policies (including combinations thereof) are beneficial. For a program pointer  $r$  that is identified as a root cause, we locate its corresponding nodes via its program location in each dynamic points-to graph, and collect (i) the number of calling contexts associated with  $r$ ,  $|cs|$ , and (ii) the sum of points-to sizes under all calling contexts,  $\sum |p_r|$ . We then count  $a_r = \frac{\sum |p_r|}{|cs|}$ , the average dynamic points-to size per calling context, to measure the accuracy of the points-to relations of  $r$  under the given context sensitivity (line 7 in Procedure 2). The context sensitivity under which  $a_r$  is the smallest is chosen for the function that contains  $r$  as the improvement suggestion (lines 10 and 11 in Procedure 2).

## 6. EVALUATION

We have conducted experiments to demonstrate the effectiveness of our approach. In this section, we present the experimental results of our analysis on various JavaScript benchmarks.

### 6.1 Experimental Setup

library	1-CFA			1-CFA + 1st-arg-sens			1-CFA + selective 1st-arg-sens			
	REAC_FUNC	AVG_TARG	HIGH_POLY	REAC_FUNC	AVG_TARG	HIGH_POLY	REAC_FUNC	AVG_TARG	HIGH_POLY	time (sec)
jQuery	312	2.2	441	206	25.0	1235	206	1.2	16	16.3
prototype.js	451	12.9	259	170	3.3	124	170	1.5	39	2.9
script.aculo.us	617	14.4	188	179	2.7	145	179	1.4	47	4.6

**Table 1:** Benchmarks I precision and performance results.

### 6.1.1 Metrics

In our evaluation, we compare the performance and precision of points-to analysis. A JavaScript points-to analysis usually constructs a call graph as well as a points-to graph. For precision on a call graph, we measure (i) **HIGH\_POLY**: the number of highly polymorphic call sites (i.e., call sites with more than 5 targets), (ii) **AVG\_TARG**: the number of targets averaged over all call sites, and (iii) **REAC\_FUNC**: the number of reachable functions. The **HIGH\_POLY** and **REAC\_FUNC** metrics were also used by Sridharan et al. [13] to evaluate the call graph construction algorithms. For precision on a points-to graph, we measure **PTS\_SIZE**: the overall points-to size (i.e., the total number of points-to set sizes over all local variables in the program; the points-to set of a variable is the number of abstract objects, represented by the allocation sites, it refers to). In addition, we measure the performance of the points-to analysis with its running time (in seconds).

### 6.1.2 Benchmarks

We use two sets of JavaScript benchmarks in our evaluation.

**Benchmarks I.** Benchmarks I consists of applications that use JavaScript libraries, generated by Sridharan et al. [13]. In this benchmark, there are 11, 5 and 1 simple web applications that invoke *jQuery*, *prototype.js* and *script.aculo.us* libraries, respectively. These libraries are among the most popular JavaScript libraries for developing real-world web applications, especially *jQuery* [14]. As reported by Sridharan et al. [13], an analysis without the specialized argument sensitivity experienced severe performance and precision problems; it required advanced static analysis techniques as well as manual code rewriting for improving precision and performance. In our experiments, we reuse these libraries with these manual transformations.

**Benchmarks II.** Benchmarks II are JavaScript applications collected by Kashyap et al. [5]. Twelve out of the 28 programs from the original benchmarks were selected for our evaluation. These programs are collected from open-source JavaScript repositories, standard JavaScript benchmarks (e.g., SunSpider<sup>3</sup>), and the Emscripten LLVM test suite<sup>4</sup>, the results of which benefit from various context-sensitive analyses [5, 15].

### 6.1.3 Experimental Design

**Root-cause localization.** We perform experiments on Benchmarks I to illustrate the accuracy of the root-cause localization algorithm. In this experiment, we use WALA’s

whole-program 1-CFA analysis as the *baseline* analysis, which experiences scalability and precision problems for the JavaScript library applications. For each of the Benchmark I programs, we perform the localization algorithm on the 1-CFA analysis, localizing a set of functions that contain the root causes. We then apply additional argument sensitivity on the first arguments (i.e., 1st-argument sensitivity [15]) of all these functions; for the rest of the functions, the 1-CFA analysis is performed (i.e., a 1-CFA and *selective* 1st-argument-sensitive analysis). We compare the precision results among the 1-CFA, the whole-program 1-CFA and 1st-argument-sensitive analysis, and the *selective* analysis using the call graph metrics for Benchmarks I. We also compare the differences in terms of analysis performance.

**Improvement suggestion.** For the programs in Benchmarks II, we apply the root-cause localization as well as improvement suggestion algorithms on the *baseline* 0-1-CFA analysis. We execute each program to obtain a program trace. A root cause function may be suggested to apply (i) a context-insensitive analysis, (ii) a single context-sensitive analysis (i.e., 1-CFA or argument sensitivity on any argument of the function<sup>5</sup>), or (iii) a combined context-sensitive analysis (e.g., 1-CFA and 1st-argument sensitivity). Based on the suggestion results, we automatically use the suggested analysis for the root cause functions; for the rest of the functions, the 0-1-CFA analysis is performed (i.e., a *auto-selective* context-sensitive analysis). We compare the performance and precision results of the *auto-selective* analysis with the 0-1-CFA analysis and the *full-sensitive* analysis (i.e., a whole-program combined context-sensitive analysis that applies 1-CFA and argument sensitivity on all arguments) for Benchmarks II to illustrate the effectiveness of the improvement suggestion.

The experimental results were obtained on a 2.5 GHz Intel Core i5 MacBook Pro with 16 GB memory running the Mac OS X 10.11 operating system.

## 6.2 Benchmarks I Results

Tables 1 and 2 show the experimental results of Benchmarks I. For each JavaScript library, the results are arithmetically averaged over all the applications that use the specific library. For example, the 291 reachable functions of *jQuery* library from the 1-CFA analysis (i.e., column 2 of the jQuery row in Table 1) is calculated by averaging the number of reachable functions the 1-CFA analysis obtained for all 11 *jQuery* applications.<sup>6</sup>

<sup>5</sup>Object sensitivity [6] applies calling contexts on the receiver argument.

<sup>6</sup>Because the applications in Benchmarks I are relatively simple programs that use the JavaScript libraries, the analysis performance and precision results of these programs are

<sup>3</sup><https://webkit.org/perf/sunspider/sunspider.html>

<sup>4</sup><http://kripken.github.io/emscripten-site/>

**Precision and performance.** Table 1 shows the precision and performance results of the analyses on Benchmarks I. Columns 2-4, 5-7, and 9-11 show the results of call graph precision metrics for the 1-CFA analysis, the whole-program combined 1-CFA and 1st-argument-sensitive analysis (i.e., *whole-combined* analysis), and the 1-CFA and *selective* 1st-argument-sensitive analysis, respectively. For each analysis, we present its number of reachable functions (i.e., columns 2, 5, and 8), average number of targets per call site (i.e., columns 3, 6 and 9), and the number of highly polymorphic call sites (i.e., columns 4, 7, and 10). In addition, column 11 show the points-to analysis time in seconds of the *selective* analysis. Given the time budget of 10 minutes, both the 1-CFA analysis and the *whole-combined* analysis failed to complete analyzing any of the programs in Benchmarks I; therefore, their precision results were calculated from the incomplete call graphs obtained after the timeout.

The **REAC\_FUNC** results of the *whole-combined* and the *selective* analyses in Table 1 are the same for all three libraries, which are significantly more precise than those of the 1-CFA analysis. Only 29% (for *script.aculo.us*) to 66% (for *jQuery*) functions considered reachable by the 1-CFA analysis are produced by the *whole-combined* and the *selective* analyses. Interestingly, the number of highly polymorphic call sites is below 50 for the *selective* analysis for all three libraries, while the 1-CFA analysis produces 188 (for *script.aculo.us*) to 441 (for *jQuery*) and the *whole-combined* analysis results in 145 (for *script.aculo.us*) to 1235 (for *jQuery*) highly polymorphic call sites. This result suggests that (i) 1-CFA analysis is imprecise to resolve the call targets in many cases, and (ii) because the *whole-combined* analysis applies 1st-argument sensitivity over all the program, it may create many calling contexts for the functions that are not identified as roots causes which may not significantly increase the analysis precision (e.g., in terms of the **REAC\_FUNC** metric). The average number of targets per call site from the *selective* analysis ranges from 1.2 (for *jQuery*) to 1.5 (for *prototype.js*), indicating the *selective* analysis precisely resolves the targets for most call sites. Although the *whole-combined* analysis reduces the **AVG\_TARG** of the 1-CFA analysis from 12.9 to 3.3 and from 14.4 to 2.7 for *prototype.js* and *script.aculo.us*, respectively, it still results in higher average number of targets per call site comparing to the *selective* analysis. For *jQuery* library, applying argument sensitivity over all the program results in significant increase of **AVG\_TARG** (i.e., on average 25 targets per call site) and **HIGH\_POLY** (i.e., 1235 highly polymorphic call sites) due to the additional calling contexts.

The last column in Table 1 shows that the *selective* analysis finishes analyzing the libraries on average between 3 seconds (for *prototype.js*) and 16 seconds (for *jQuery*). Due to the fact that the 1-CFA analysis could not finish analyzing any of these libraries under 10 minutes, we claim that the *selective* analysis has significantly better performance because of the precision gained by applying 1st-argument sensitivity on the root cause functions. The *whole-combined* analysis also could not complete within the 10-minute time budget. Applying 1st-argument sensitivity over all the program generates too many calling contexts for the propagation system to converge. In summary, the results in Table 1 suggest that (i) our root-cause localization algorithm is capable of locating the functions where the 1-CFA analysis experiences sig-

dominated by the underlying libraries. Therefore, we report the average results based on the corresponding libraries.

library	no. of localized functions	no. of evaluations	slope
jQuery	2	25000	3.7
prototype.js	4	72000	5.1
script.aculo.us	4	96000	4.9

**Table 2:** Root-cause localization results of Benchmarks I.

nificant precision and performance loss, and (ii) because the root causes are highly condensed in these JavaScript libraries (see more discussions below), applying the 1st-argument-sensitive analysis only on the localized functions achieved a much better balance between precision and performance comparing to the *whole-combined* analysis.

**Root-cause localization characteristics.** Table 2 shows additional information that characterizes the results of the root-cause localization algorithm. For the experiments on Benchmarks I, we paused the 1-CFA analysis every 1000 evaluations to decide if the root-cause localization algorithm should be performed. Columns 2, 3, and 4 present the number of functions identified as root causes, the number of evaluations until performing root-cause localization, and the slope of the last 1000 evaluations (i.e., the increase in the total number of points-to relations divided by 1000), respectively.

Our algorithm identifies 2, 4 and 4 functions as root causes for *jQuery*, *prototype.js*, and *script.aculo.us*, respectively. Comparing to the number of reachable functions computed by any analysis in Table 1 (i.e., more than 150 functions), very small fractions of these functions were identified as root causes. This result as well as the good precision and performance of the *selective* analysis support our intuition that a small number of complex constructs in the programs may contribute to significant loss of analysis performance and precision if not handled accurately. Therefore, it is useful for our automated localization algorithm to pinpoint these root causes, as shown.

Column 4 shows that the slopes of the last 1000 evaluations range from 3.7 (for *jQuery*) to 5.1 (for *prototype.js*). For example for *prototype.js*, the large slope indicates that the precision of the 1-CFA analysis significantly decreases during this period (i.e., about 5 new points-to relations per constraint), while the slope of the simple linear regression for all previous evaluations is 0.9. This result suggests that our heuristics accurately decide when to perform the root-cause localization for JavaScript libraries.

### 6.3 Benchmarks II Results

Figures 6 and 7 show the precision and performance results of Benchmarks II, respectively. Because the 0-1-CFA analysis finishes analyzing all 12 programs in Benchmarks II within the time budget of 10 minutes, the root-cause localization was performed after the 0-1-CFA points-to analysis completes on each program.

Figure 6 presents the **PTS\_SIZE** precision improvement of the *full-sensitive* analysis (i.e., grey bars) and the *auto-selective* analysis (i.e., patterned bars) over the 0-1-CFA analysis. Because the *full-sensitive* analysis applies 1-CFA and argument sensitivity of all arguments over all the func-



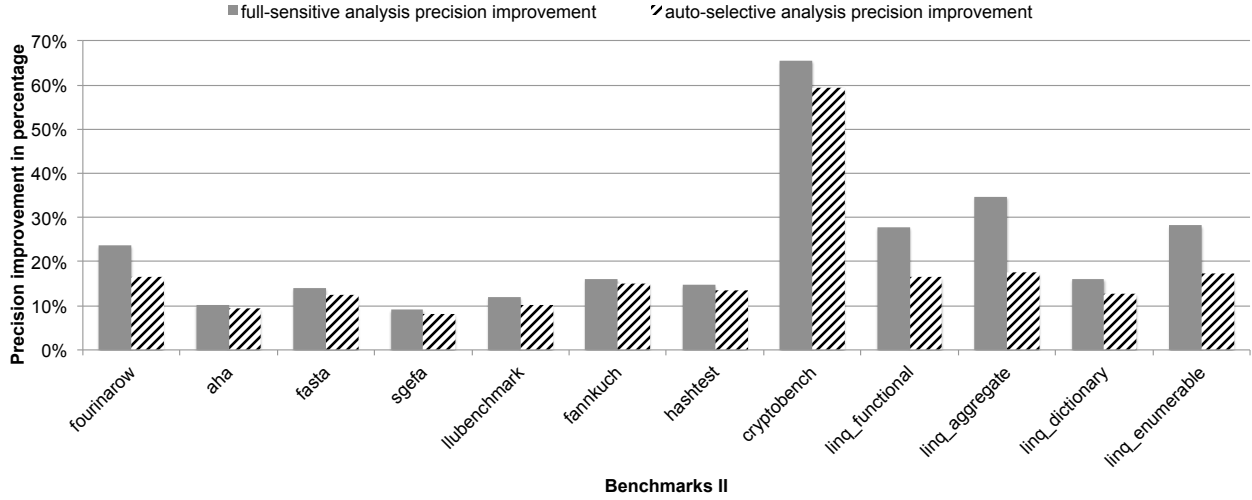


Figure 6: Benchmarks II precision results.

tions, its results are at least as precise as the *selective* analysis. In Figure 6, the y axis shows the precision improvement of the corresponding analysis Y over the 0-1-CFA analysis 0-1-CFA, calculated as follows

$$\text{IMP}_Y = \frac{\text{PTS\_SIZE}_{0-1\text{-CFA}} - \text{PTS\_SIZE}_Y}{\text{PTS\_SIZE}_{0-1\text{-CFA}}} \times 100\%$$

Therefore,  $\text{IMP}_Y$  measures analysis Y's precision in terms of removing the false positives from the 0-1-CFA analysis results. In Figure 6, the *whole-combined* analysis improves the PTS\_SIZE precision over the 0-1-CFA analysis by between 9% (for *sgefa*) to 65% (for *cryptobench*). For all but five programs (i.e., *fourinarow*, *cryptobench*, *linq\_functional*, *linq\_aggregate* and *linq\_enumerable*), the differences of the precision improvement percentages are within 3.5% between the *full-sensitive* and the *auto-selective* analyses, indicating that the *auto-selective* analysis produces similar results to the *full-sensitive* analysis for most programs. The results of the *linq\_functional* program exhibit the largest difference in terms of precision improvement (i.e., 17%) between the *full-sensitive* 35%) and the *auto-selective* (18%) analyses. Nevertheless, more than 50% of the false positives that are produced by the 0-1-CFA analysis, but not the *full-sensitive* analysis, are absent from the *auto-selective* analysis results for this program.

Figure 7 presents the performance results of the *full-sensitive* (i.e., grey bars) and the *auto-selective* (i.e., patterned bars) analyses comparing to the 0-1-CFA analysis performance. The y axis shows the overhead of the corresponding analysis Y's time cost comparing to the 0-1-CFA analysis time (i.e.,  $\frac{\text{TIME}_Y}{\text{TIME}_{0-1\text{-CFA}}}$ ). The performance of an analysis on each benchmark program was obtained by averaging over 30 repeated executions.

In Figure 7, the *full-sensitive* analysis performs significantly worse than the 0-1-CFA analysis for all the benchmark programs. It could not finish analyzing *linq\_aggregate* and *linq\_enumerable* under the time budget of 10 minutes; therefore, the incomplete points-to results of these two programs were obtained after the timeout for comparison. The *full-sensitive* analysis is at least two orders of magnitude slower for another three program (i.e., *linq\_functional*, *aha*,

and *linq\_dictionary*) and is between 23 (for *fannkuch*) and 60 (for *llubenchmark*) times slower for six programs than the 0-1-CFA analysis. For example, it takes less than one second for the 0-1-CFA analysis to finish analyzing *linq\_functional*, while the *full-sensitive* analysis needs almost 7 minutes to complete analyzing the same program. Despite of the fact that the *full-sensitive* analysis often results in relatively significant precision improvement (e.g., 28% for *linq\_functional*), the performance issues outweigh the benefits of this whole-program combined context-sensitive analysis in many cases.

On the other hand, the *auto-selective* analysis is capable of analyzing the benchmarks under the same order of magnitude as the 0-1-CFA analysis for all but three programs (i.e., *fourinarow*, *aha* and *linq\_dictionary*). For example, the *auto-selective* analysis improves the precision over the 0-1-CFA analysis by 59% for *cryptobench* and it also finishes analyzing this program almost as fast as the 0-1-CFA analysis. In most cases, the *auto-selective* analysis results in a much better balance between performance and precision than the *full-sensitive* analysis (e.g., the *full-sensitive* analysis is 65% more precise but performs 5 times slower than the 0-1-CFA analysis for *cryptobench*). For three programs the *auto-selective* analysis performs more than 10 times slower than the 0-1-CFA analysis (i.e., 10, 15 and 143 times slower for *aha*, *fourinarow* and *linq\_dictionary*, respectively), the *auto-selective* analysis is still an order of magnitude faster than the *full-sensitive* analysis for *aha* and about 3 times faster than the *full-sensitive* analysis for *fourinarow*. An outlier is the performance of the *auto-selective* analysis for *linq\_dictionary*, which is similar to the *full-sensitive* analysis in performance for *linq\_dictionary*, indicating the *auto-selective* analysis may have applied expensive context sensitivity on a set of functions that result in performance overhead. Nevertheless, the *auto-selective* analysis has achieved significantly better performance than the *full-sensitive* analysis for most of the programs in Benchmarks II.

Our localization algorithm identifies between 6% (for *fourinarow*) to 27% (for *linq\_aggregate*) of the functions as the sources of precision loss, with an average of 13% of the functions over all the programs in Benchmarks II, a relatively small fraction. Our improvement suggestion algo-




Figure 7: Benchmarks II performance results.

rithm recommends 72%, 25%, and 3% of all the root cause functions in Benchmarks II to apply a combined context-sensitive analysis, a single context-sensitive analysis and a context-insensitive analysis, respectively.

In summary, the *auto-selective* analysis obtains similar precision results to the *full-sensitive* analysis that improves the precision of the 0-1-CFA analysis; moreover, the *auto-selective* analysis’ performance is significantly better than the whole-program combined fully context-sensitive analysis. This result suggests that our localization algorithm can accurately identify the small fraction of functions as root causes and our improvement suggestion algorithm can suggest the appropriate context sensitivity that benefit the results both in performance and precision for Benchmarks II.

## 7. RELATED WORK

To the best of our knowledge, we present the first work that focuses on automating the process of localizing the root causes when a JavaScript static analysis produces impractical results on a program. Nevertheless, our work is related to (i) recent static context-sensitive analyses that unveiled difficult JavaScript program constructs, and (ii) static analysis that identifies the causes of precision loss for refinement.

### 7.1 JavaScript Context-sensitive Analysis

Sridharan et al. identify correlated dynamic property accesses (e.g.,  $x[p] = y[p]$ ) as a hard-to-analyze JavaScript code pattern and present a specialized argument-sensitive analysis along with program transformation that dramatically improves analysis scalability and precision on JavaScript libraries [13]. This work motivated us to design an automatic localization algorithm that identifies such program constructs causing the analysis imprecision. We also have reused the library benchmarks collected by Sridharan et al. [13] as Benchmarks I in the evaluation.

Esben et al. infer determinacy information (i.e., a variable and expression always has the same value at a given program point [8]) and improve static analysis precision on JavaScript libraries via various techniques [2]. For example, argument sensitivity is selectively applied for the arguments whose abstract value is a concrete string or a single object address.

This specialized context sensitivity is proposed as a result of manually inspecting imprecise portions of an incomplete call graph obtained after a fixed number of evaluations. In our work, root causes of the analysis imprecision are automatically localized and a user is provided with focused program constructs for inspection.

Wei and Ryder present a two-staged context-sensitive analysis for JavaScript that selectively applies specific context sensitivity on the function level [15]. Heuristics are used to decide the context sensitivity based on the function characteristics extracted from the results of a pre-analysis. Our improvement suggestion may also result in applying different context sensitivity on the root cause functions. However, our approach does not require a pre-analysis that finishes analyzing the target program and our goal is to apply the specialized context sensitivity only on the root cause functions that may significantly improve the overall analysis scalability and precision.

Park and Ryu present another static analysis of JavaScript that improves precision via loop sensitivity [7]. The authors identify one root cause of scalability problems with JavaScript analysis as the combination of imprecise results in loops and in dynamic property accesses. Their analysis improves precision in loops by distinguishing each iteration of a loop with different contexts based on the analysis results of loop conditional expressions. Our approach systematically assists in the process of locating such code patterns as root causes of analysis scalability problems.

### 7.2 Refinement-based Analysis

Smaragdakis et al. present introspective analysis that aims to improve the performance of a context-sensitive analysis for Java [11]. Introspective analysis uses heuristics to decide whether to refine an allocation site or a call site with context sensitivity, based on the metrics computed from context-insensitive points-to results. The heuristics focus on reasoning about the cost of applying additional context sensitivity. Instead, we focus on identifying the constructs that originate significant loss of performance and/or precision of an analysis; therefore, applying more accurate but expensive analysis techniques only on these constructs may result

in the improvement of the overall analysis performance and precision.

Sridharan and Bodík present a refinement-based points-to analysis for Java that refines sensitivity for heap accesses and method calls [12]. The analysis is demand-driven and client-driven in that it focuses on refining the analysis of relevant code. We focus on improving the overall points-to analysis precision in terms of all possible queries for points-to sets of program variables, instead of precisely answering a specific query.

Guyer and Lin present a client-driven analysis for C that automatically adjusts its precision in response to the needs of client analyses [4]. This client-driven analysis monitors polluting assignments (i.e., the program points that result in inaccuracy in the analysis) and tunes context as well as flow sensitivity to improve precision. Our heuristics to locate the root causes consider not only the inaccurate results at a program point but also its impact on the overall points-to analysis precision by tracking the labels in the propagation system.

## 8. CONCLUSIONS

Discovering the causes of a static analysis being impractical is critical for developing novel automated software tools. In this paper, we have designed the first systematic support for this time-consuming process, presenting automatic root-cause localization and improvement suggestion. To diagnose a JavaScript points-to analysis, we instrument the propagation system with labels that keep track of this history of points-to propagation. This information is then used to identify the pointers whose imprecision have large impact on the overall analysis performance and/or precision. The improvement suggestion algorithm takes the dynamic information to simulate the benefits of various context-sensitive analyses to improve the precision on the root causes via context sensitivity. We have performed evaluation on two sets of benchmarks. The results on library applications suggest that the localization algorithm is capable of identifying a small set of root causes that significantly affect the analysis performance and precision. Applying context sensitivity specifically on these root cause functions resolves the scalability problems that occur on both the imprecise analysis and the whole-program combined context-sensitive analysis. The analysis that automatically applies the suggested context sensitivity on the root cause functions achieves a better balance between precision and performance for most programs in Benchmarks II, demonstrating the effectiveness of the improvement suggestion algorithm.

In the future, we would like to improve the usability of our approach with an interactive user interface; therefore, it becomes more useful for both analysis designer and end user for developing new program analysis techniques and building software tools. We plan to further investigate the currently preliminary improvement suggestion for including other advanced analysis techniques (e.g., program transformation) in the recommendations.

## 9. REFERENCES

- [1] L. O. Andersen. *Program Analysis and Specialization for the C Programming Language*. PhD thesis, University of Copenhagen, 1994.
- [2] E. Andreassen and A. Møller. Determinacy in static analysis for jQuery. In *Proceedings of the 2014 ACM International Conference on Object Oriented Programming Systems Languages & Applications, OOPSLA '14*, pages 17–31, New York, NY, USA, 2014. ACM.
- [3] R. Cytron, J. Ferrante, B. K. Rosen, M. N. Wegman, and F. K. Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Trans. Program. Lang. Syst.*, 13(4):451–490, Oct. 1991.
- [4] S. Z. Guyer and C. Lin. Client-driven pointer analysis. In *Proceedings of the 10th International Conference on Static Analysis, SAS'03*, pages 214–236, 2003.
- [5] V. Kashyap, K. Dewey, E. A. Kuefner, J. Wagner, K. Gibbons, J. Sarracino, B. Wiedermann, and B. Hardekopf. Jsai: A static analysis platform for JavaScript. In *Proceedings of the 22nd ACM SIGSOFT International Symposium on Foundations of Software Engineering, FSE 2014*, pages 121–132, New York, NY, USA, 2014. ACM.
- [6] A. Milanova, A. Rountev, and B. G. Ryder. Parameterized object sensitivity for points-to analysis for Java. *ACM Trans. Softw. Eng. Methodol.*, 14(1):1–41, Jan. 2005.
- [7] C. Park and S. Ryu. Scalable and precise static analysis of JavaScript applications via loop-sensitivity. In *29th European Conference on Object-Oriented Programming, ECOOP 2015, July 5-10, 2015, Prague, Czech Republic*, pages 735–756, 2015.
- [8] M. Schäfer, M. Sridharan, J. Dolby, and F. Tip. Dynamic determinacy analysis. In *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI '13*, pages 165–174, New York, NY, USA, 2013. ACM.
- [9] M. Sharir and A. Pnueli. Two approaches to interprocedural data flow analysis. *Program Flow Analysis: Theory and Applications*, pages 189–234, 1981.
- [10] O. G. Shivers. *Control-flow Analysis of Higher-order Languages of Taming Lambda*. PhD thesis, Carnegie Mellon University, 1991.
- [11] Y. Smaragdakis, G. Kastrinis, and G. Balatsouras. Introspective analysis: Context-sensitivity, across the board. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI '14*, pages 485–495, 2014.
- [12] M. Sridharan and R. Bodík. Refinement-based context-sensitive points-to analysis for Java. In *Proceedings of the 2006 ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI '06*, pages 387–400, 2006.
- [13] M. Sridharan, J. Dolby, S. Chandra, M. Schäfer, and F. Tip. Correlation tracking for points-to analysis of JavaScript. In *Proceedings of the 26th European Conference on Object-Oriented Programming, ECOOP'12*, pages 435–458, 2012.
- [14] W<sup>3</sup>Techns. W<sup>3</sup>Techns web technologies surveys: usage of JavaScript libraries for websites. <http://w3techns.com/technologies/overview/javascript.library/all>, 2016.
- [15] S. Wei and B. G. Ryder. Adaptive context-sensitive analysis for JavaScript. In *29th European Conference on Object-Oriented Programming, ECOOP 2015, July 5-10, 2015, Prague, Czech Republic*, pages 712–734,

2015.